

TIME TO READMISSSION IN DIABETIC PATIENTS USING SURVIVAL ANALYSIS

Created by: AJIROLA AMUDAT

```
In [1]: # Import Required Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from lifelines import KaplanMeierFitter, CoxPHFitter
from lifelines.statistics import logrank_test
from lifelines.statistics import multivariate_logrank_test
```

```
C:\Users\HELLO\anaconda3\Lib\site-packages\pandas\core\arrays\masked.py:60: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).
  from pandas.core import (
```

```
In [2]: # To Load Dataset
data = pd.read_csv("diabetic_data.csv")
```

```
In [3]: # Show the first 10 rows of the dataset
data.head(10)
```

Out [3]:

	encounter_id	patient_nbr	race	gender	age	weight	admission_type_id	discharge_disposition_id	admission_source_id	time_in_hospital	...	citoglipton
0	2278392	8222157	Caucasian	Female	[0-10)	?	6	25	1	1	...	No
1	149190	55629189	Caucasian	Female	[10-20)	?	1	1	7	3	...	No
2	64410	86047875	AfricanAmerican	Female	[20-30)	?	1	1	7	2	...	No
3	500364	82442376	Caucasian	Male	[30-40)	?	1	1	7	2	...	No
4	16680	42519267	Caucasian	Male	[40-50)	?	1	1	7	1	...	No
5	35754	82637451	Caucasian	Male	[50-60)	?	2	1	2	3	...	No
6	55842	84259809	Caucasian	Male	[60-70)	?	3	1	2	4	...	No
7	63768	114882984	Caucasian	Male	[70-80)	?	1	1	7	5	...	No
8	12522	48330783	Caucasian	Female	[80-90)	?	2	1	4	13	...	No
9	15738	63555939	Caucasian	Female	[90-100)	?	3	3	4	12	...	No

10 rows × 50 columns

DATA CLEANING

```
In [4]: # To know the dataset columns
data.columns
```

Out [4]: Index(['encounter_id', 'patient_nbr', 'race', 'gender', 'age', 'weight', 'admission_type_id', 'discharge_disposition_id', 'admission_source_id', 'time_in_hospital', 'payer_code', 'medical_specialty', 'num_lab_procedures', 'num_procedures', 'num_medications', 'number_outpatient', 'number_emergency', 'number_inpatient', 'diag_1', 'diag_2', 'diag_3', 'number_diagnoses', 'max_glu_serum', 'A1Cresult', 'metformin', 'repaglinide', 'nateglinide', 'chlorpropamide', 'glimepiride', 'acetohexamide', 'glipizide', 'glyburide', 'tolbutamide', 'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone', 'tolazamide', 'examide', 'citoglipton', 'insulin', 'glyburide-metformin', 'glipizide-metformin', 'glimepiride-pioglitazone', 'metformin-rosiglitazone', 'metformin-pioglitazone', 'change', 'diabetesMed', 'readmitted'], dtype='object')

```
In [5]: # To know the shape or size of the dataset
data.shape
```

Out [5]: (101766, 50)

```
In [6]: # To know the data types and get familiarise with the data to know how to work with them.
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101766 entries, 0 to 101765
Data columns (total 50 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   encounter_id                          101766 non-null  int64
 1   patient_nbr                           101766 non-null  int64
 2   race                                  101766 non-null  object
 3   gender                                101766 non-null  object
 4   age                                   101766 non-null  object
 5   weight                                101766 non-null  object
 6   admission_type_id                     101766 non-null  int64
 7   discharge_disposition_id              101766 non-null  int64
 8   admission_source_id                   101766 non-null  int64
 9   time_in_hospital                      101766 non-null  int64
10   payer_code                            101766 non-null  object
11   medical_specialty                     101766 non-null  object
12   num_lab_procedures                    101766 non-null  int64
13   num_procedures                        101766 non-null  int64
14   num_medications                       101766 non-null  int64
15   number_outpatient                     101766 non-null  int64
16   number_emergency                      101766 non-null  int64
17   number_inpatient                      101766 non-null  int64
18   diag_1                                101766 non-null  object
19   diag_2                                101766 non-null  object
20   diag_3                                101766 non-null  object
21   number_diagnoses                      101766 non-null  int64
22   max_glu_serum                         5346 non-null   object
23   A1Cresult                             17018 non-null  object
24   metformin                             101766 non-null  object
25   repaglinide                           101766 non-null  object
26   nateglinide                           101766 non-null  object
```

```

27 chlorpropamide 101766 non-null object
28 glimepiride 101766 non-null object
29 acetohexamide 101766 non-null object
30 glipizide 101766 non-null object
31 glyburide 101766 non-null object
32 tolbutamide 101766 non-null object
33 pioglitazone 101766 non-null object
34 rosiglitazone 101766 non-null object
35 acarbose 101766 non-null object
36 miglitol 101766 non-null object
37 troglitazone 101766 non-null object
38 tolazamide 101766 non-null object
39 examide 101766 non-null object
40 citoglipton 101766 non-null object
41 insulin 101766 non-null object
42 glyburide-metformin 101766 non-null object
43 glipizide-metformin 101766 non-null object
44 glimepiride-pioglitazone 101766 non-null object
45 metformin-rosiglitazone 101766 non-null object
46 metformin-pioglitazone 101766 non-null object
47 change 101766 non-null object
48 diabetesMed 101766 non-null object
49 readmitted 101766 non-null object
dtypes: int64(13), object(37)
memory usage: 38.8+ MB

```

```
In [7]: # To check for missing value
data.isnull().sum()
```

```

Out [7]: encounter_id      0
patient_nbr      0
race             0
gender           0
age              0
weight           0
admission_type_id 0
discharge_disposition_id 0
admission_source_id 0
time_in_hospital 0
payer_code       0
medical_specialty 0
num_lab_procedures 0
num_procedures   0
num_medications  0
number_outpatient 0
number_emergency 0
number_inpatient 0
diag_1           0
diag_2           0
diag_3           0
number_diagnoses 0
max_glu_serum    96420
A1CResult        84748
metformin        0
repaglinide      0
nateglinide      0
chlorpropamide   0
glimepiride      0
acetohexamide    0
glipizide        0
glyburide        0
tolbutamide      0
pioglitazone     0
rosiglitazone    0
acarbose         0
miglitol         0
troglitazone     0
tolazamide       0
examide          0
citoglipton      0
insulin          0
glyburide-metformin 0
glipizide-metformin 0
glimepiride-pioglitazone 0
metformin-rosiglitazone 0
metformin-pioglitazone 0
change           0
diabetesMed      0
readmitted       0
dtype: int64

```

```
In [8]: # To Keep necessary variables
data = data[['age', 'time_in_hospital', 'num_lab_procedures',
            'num_medications', 'number_diagnoses', 'insulin', 'readmitted']]
```

```
In [9]: # To know the sum of duplicates in the dataset
duplicate_count = data.duplicated().sum()
duplicate_count
```

```
Out [9]: 3877
```

```
In [10]: # To remove duplicate
data.drop_duplicates()
```

```

Out [10]:
   age  time_in_hospital  num_lab_procedures  num_medications  number_diagnoses  insulin  readmitted
0  [0-10]  1             41                      1                1             No      NO
1  [10-20]  3             59                      18                9             Up      >30
2  [20-30]  2             11                      13                6             No      NO
3  [30-40]  2             44                      16                7             Up      NO
4  [40-50]  1             51                      8                 5             Steady  NO
...  ...  ...             ...                      ...                ...             ...      ...
101761  [70-80]  3             51                      16                9             Down  >30
101762  [80-90]  5             33                      18                9             Steady  NO
101763  [70-80]  1             53                      9                13            Down  NO
101764  [80-90]  10            45                      21                9             Up      NO
101765  [70-80]  6             13                      3                 9             No      NO

```

97889 rows × 7 columns

```
In [11]: # To Remove rows with missing age
data = data[data['age'] != '?']
```

```
In [12]: # To Keep valid insulin types
data = data[data['insulin'].isin(['Up', 'Down', 'Steady'])]
```

```
In [13]: # To Define event (readmitted <30 days = 1) and duration (time_in_hospital)
data['event'] = data['readmitted'].apply(lambda x: 1 if x == '<30' else 0)
data['duration'] = data['time_in_hospital']
```

```
In [14]: # To create ordered age categories
age_order = ['[0-10]', '[10-20]', '[20-30]', '[30-40]', '[40-50]', '[50-60]',
```

```
data['age'] = pd.Categorical(data['age'], categories=age_order, ordered=True)
```

```
In [15]: # To create medication groupings
data['med_group'] = pd.cut(data['num_medications'],
                           bins=[0, 10, 20, 40, data['num_medications'].max()],
                           labels=['Low', 'Moderate', 'High', 'Very High'],
                           include_lowest=True)
```

```
In [16]: # To show the first 10 rows of the dataset
data.head(10)
```

Out [16]:

	age	time_in_hospital	num_lab_procedures	num_medications	number_diagnoses	insulin	readmitted	event	duration	med_group
1	[10-20)	3	59	18	9	Up	>30	0	3	Moderate
3	[30-40)	2	44	16	7	Up	NO	0	2	Moderate
4	[40-50)	1	51	8	5	Steady	NO	0	1	Low
5	[50-60)	3	31	16	9	Steady	>30	0	3	Moderate
6	[60-70)	4	70	21	7	Steady	NO	0	4	High
8	[80-90)	13	68	28	8	Steady	NO	0	13	High
9	[90-100)	12	33	18	8	Steady	NO	0	12	Moderate
10	[40-50)	9	47	17	9	Steady	>30	0	9	Moderate
11	[60-70)	7	62	11	7	Steady	<30	1	7	Moderate
12	[40-50)	7	60	15	8	Down	<30	1	7	Moderate

```
In [17]: # To know the shape or size of the dataset
data.shape
```

Out [17]: (54383, 10)

EXPLORATORY DATA ANALYSIS (EDA)

```
In [18]: # To get the summary statistics of the data
data.describe()
```

Out [18]:

	time_in_hospital	num_lab_procedures	num_medications	number_diagnoses	event	duration
count	54383.000000	54383.000000	54383.000000	54383.000000	54383.000000	54383.000000
mean	4.671956	44.811191	17.601144	7.566188	0.121380	4.671956
std	3.052438	19.841404	8.606966	1.917152	0.326571	3.052438
min	1.000000	1.000000	1.000000	1.000000	0.000000	1.000000
25%	2.000000	33.000000	12.000000	6.000000	0.000000	2.000000
50%	4.000000	46.000000	16.000000	9.000000	0.000000	4.000000
75%	6.000000	59.000000	22.000000	9.000000	0.000000	6.000000
max	14.000000	132.000000	81.000000	16.000000	1.000000	14.000000

KAPLAN-MEIR SURVIVAL CURVE

Overall Kaplan-Meier Estimator


```
In [19]: # Overall Kaplan-Meier Estimator
kmf = KaplanMeierFitter()
```

```
In [20]: # To fit KM model
kmf.fit(data['duration'], event_observed=data['event'])
```

Out [20]: <lifelines.KaplanMeierFitter:"KM_estimate", fitted with 54383 total observations, 47782 right-censored observations>

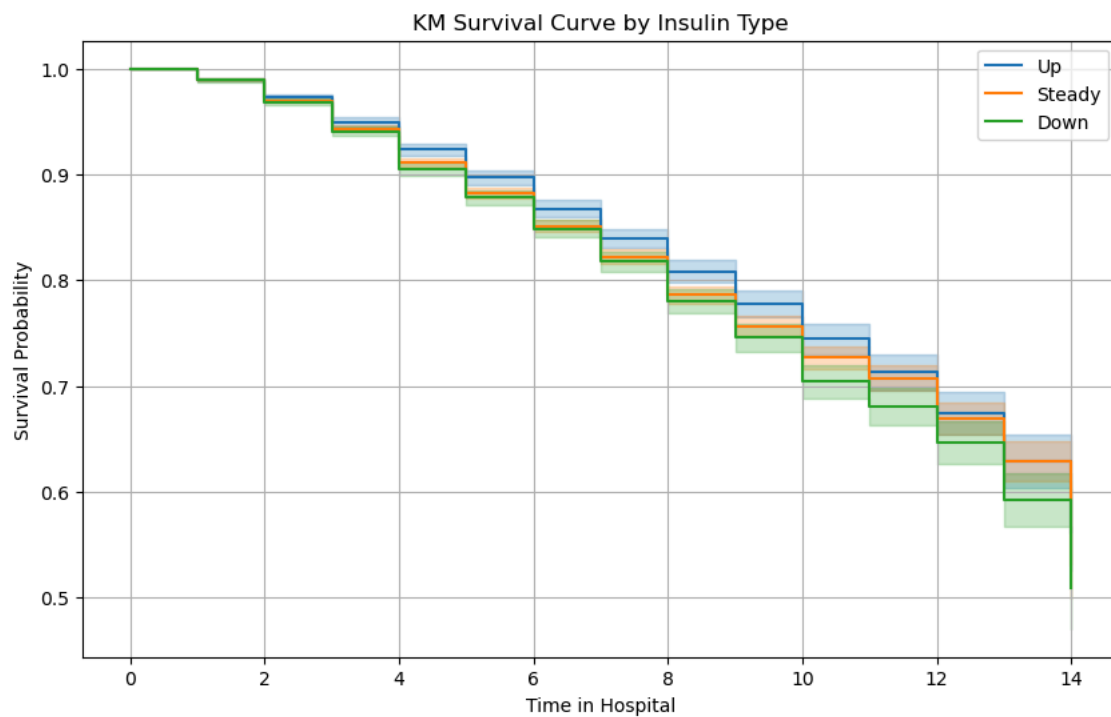
```
In [21]: # To Plot KM survival function
kmf.plot_survival_function()
plt.title("Kaplan-Meier Survival Curve")
plt.xlabel("Time in Hospital (Days)")
plt.ylabel("Survival Probability")
plt.grid(True)
plt.show()
```

Kaplan-Meier Survival Curve



Kaplan-Meier Curve by Insulin Type


```
In [22]: # To plot Kaplan-Meier curve by insulin type
kmf = KaplanMeierFitter()
plt.figure(figsize=(10, 6))
for grp in data['insulin'].dropna().unique():
    grp_data = data[data['insulin'] == grp]
    kmf.fit(grp_data['duration'], grp_data['event'], label=grp)
    kmf.plot_survival_function()
plt.title("KM Survival Curve by Insulin Type")
plt.xlabel("Time in Hospital")
plt.ylabel("Survival Probability")
plt.legend()
plt.grid(True)
plt.show()
```



Kaplan-Meier curve by Age group

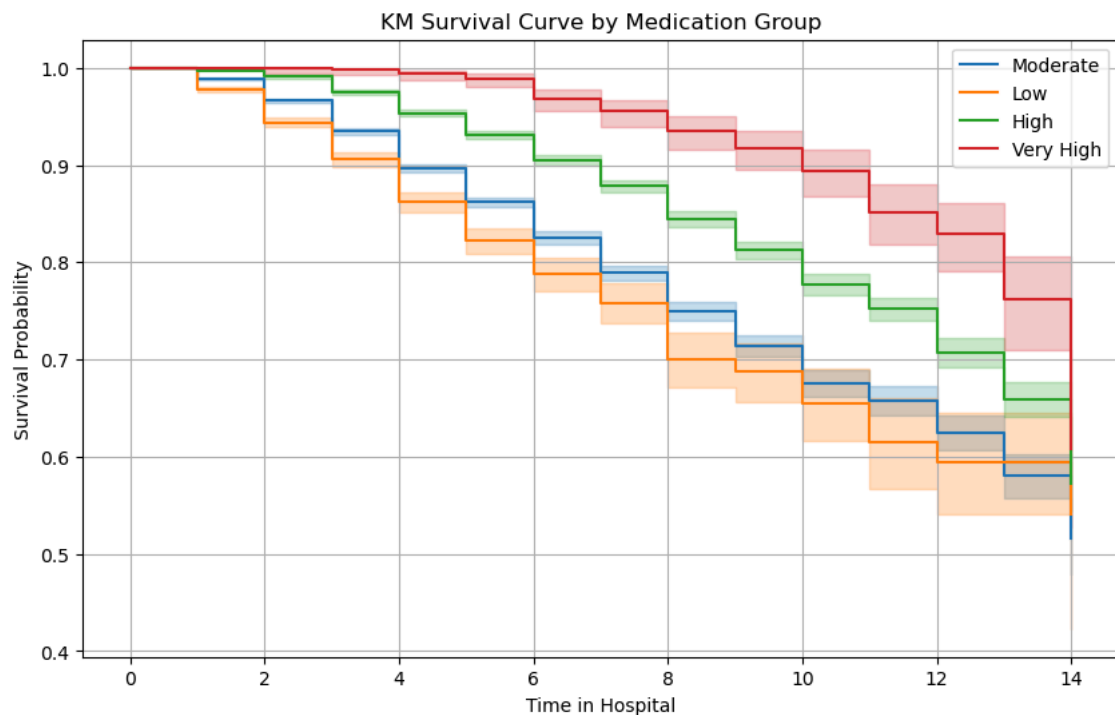
```
In [23]: # To plot Kaplan-Meier curve by Age group
plt.figure(figsize=(12, 6))
for grp in data['age'].cat.categories:
    grp_data = data[data['age'] == grp]
    if not grp_data.empty:
        kmf.fit(grp_data['duration'], grp_data['event'], label=str(grp))
        kmf.plot_survival_function()
plt.title("KM Survival Curve by Age Group")
plt.xlabel("Time in Hospital")
plt.ylabel("Survival Probability")
plt.legend()
plt.grid(True)
plt.show()
```

KM Survival Curve by Age Group



Kaplan-Meier curve by Medication group

```
In [24]: # To plot Kaplan-Meier curve by medication group
plt.figure(figsize=(10, 6))
for grp in data['med_group'].dropna().unique():
    grp_data = data[data['med_group'] == grp]
    kmf.fit(grp_data['duration'], grp_data['event'], label=str(grp))
    kmf.plot_survival_function()
plt.title("KM Survival Curve by Medication Group")
plt.xlabel("Time in Hospital")
plt.ylabel("Survival Probability")
plt.legend()
plt.grid(True)
plt.show()
```



LOG RANK TEST

Log-rank tests to compare survival curves across groups

```
In [25]: # To compare survival curves across groups
print(" Log-Rank Test: Insulin")
results_insulin = print(multivariate_logrank_test(data['duration'], data['insulin'], data['event']).summary)
print(results_insulin)
print(" Log-Rank Test: Age")
results_age = print(multivariate_logrank_test(data['duration'], data['age'], data['event']).summary)
print(results_age)
print(" Log-Rank Test: Medication Group")
results_medication = print(multivariate_logrank_test(data['duration'], data['med_group'], data['event']).summary)
print(results_medication)
```

```
Log-Rank Test: Insulin
test_statistic      p      -log2(p)
0      19.30635      0.000064      13.926588
None
Log-Rank Test: Age
test_statistic      p      -log2(p)
0      124.428671      1.653295e-22      72.357074
None
Log-Rank Test: Medication Group
test_statistic      p      -log2(p)
```

Cox Proportional Hazards Model

```
In [26]: # To convert categorical variables to dummy variables
data_encoded = pd.get_dummies(data[['age', 'insulin', 'med_group', 'duration', 'event']],
                               columns=['age', 'insulin', 'med_group'], drop_first=True)

In [27]: # To fit Cox Model
cph = CoxPHFitter()
cph.fit(data_encoded, duration_col='duration', event_col='event')
cph.print_summary()
```

model	lifelines.CoxPHFitter												
duration col	'duration'												
event col	'event'												
baseline estimation	breslow												
number of observations	54383												
number of events observed	6601												
partial log-likelihood	-65285.39												
time fit was run	2025-05-26 11:08:57 UTC												
	coef	exp(coef)	se(coef)	coef lower 95%	coef upper 95%	exp(coef) lower 95%	exp(coef) upper 95%	cmp to	z	p	-log2(p)		
age_[10-20]	0.71	2.04	0.60	-0.46	1.89	0.63	6.63	0.00	1.19	0.24	2.09		
age_[20-30]	1.65	5.21	0.58	0.51	2.79	1.66	16.31	0.00	2.84	<0.005	7.77		
age_[30-40]	1.35	3.85	0.58	0.21	2.49	1.23	12.01	0.00	2.32	0.02	5.62		
age_[40-50]	1.27	3.54	0.58	0.13	2.40	1.14	11.03	0.00	2.18	0.03	5.11		
age_[50-60]	1.14	3.13	0.58	0.01	2.28	1.01	9.75	0.00	1.97	0.05	4.37		
age_[60-70]	1.19	3.29	0.58	0.06	2.32	1.06	10.23	0.00	2.06	0.04	4.66		
age_[70-80]	1.15	3.15	0.58	0.01	2.28	1.01	9.78	0.00	1.98	0.05	4.39		
age_[80-90]	1.12	3.08	0.58	-0.01	2.26	0.99	9.56	0.00	1.94	0.05	4.26		
age_[90-100]	1.08	2.96	0.58	-0.06	2.23	0.94	9.27	0.00	1.86	0.06	3.99		
insulin_Steady	-0.14	0.87	0.03	-0.20	-0.08	0.82	0.92	0.00	-4.72	<0.005	18.72		
insulin_Up	-0.14	0.87	0.04	-0.21	-0.07	0.81	0.93	0.00	-3.97	<0.005	13.76		
med_group_Moderate	-0.34	0.71	0.04	-0.42	-0.27	0.66	0.76	0.00	-9.29	<0.005	65.81		
med_group_High	-0.90	0.41	0.04	-0.98	-0.82	0.38	0.44	0.00	-21.70	<0.005	344.38		
med_group_Very High	-1.47	0.23	0.10	-1.66	-1.28	0.19	0.28	0.00	-15.36	<0.005	174.36		
Concordance	0.63												
Partial AIC	130598.77												
log-likelihood ratio test	836.15 on 14 df												
-log2(p) of ll-ratio test	560.38												

The cox model assumptions

```
In [28]: # To test the cox model assumptions
cph.fit(data_encoded, duration_col='duration', event_col='event')
cph.check_assumptions(data_encoded, p_value_threshold=0.05)
```

The ``p_value_threshold`` is set at 0.05. Even under the null hypothesis of no violations, some covariates will be below the threshold by chance. This is compounded when there are many covariates. Similarly, when there are lots of observations, even minor deviances from the proportional hazard assumption will be flagged.

With that in mind, it's best to use a combination of statistical tests and visual tests to determine the most serious violations. Produce visual plots using ``check_assumptions(..., show_plots=True)`` and looking for non-constant lines. See link [A] below for a full example.

null_distribution	chi squared			
degrees_of_freedom	1			
model	<lifelines.CoxPHFitter: fitted with 54383 tota...			
test_name	proportional_hazard_test			
	test_statistic	p	-log2(p)	
age_[10-20]	km	0.42	0.52	0.95
	rank	0.27	0.60	0.74
age_[20-30]	km	0.39	0.53	0.90
	rank	0.29	0.59	0.76
age_[30-40]	km	0.06	0.80	0.32
	rank	0.02	0.89	0.16
age_[40-50]	km	0.09	0.77	0.38
	rank	0.03	0.86	0.22
age_[50-60]	km	0.04	0.84	0.26
	rank	0.00	0.96	0.07
age_[60-70]	km	0.07	0.80	0.33
	rank	0.01	0.94	0.09
age_[70-80]	km	0.03	0.87	0.20
	rank			

		test_statistic	p	-log2(p)
	rank	0.00	1.00	0.01
age_[80-90)	km	0.02	0.89	0.17
	rank	0.02	0.89	0.18
age_[90-100)	km	0.08	0.77	0.37
	rank	0.01	0.94	0.08
insulin_Steady	km	0.13	0.72	0.48
	rank	0.28	0.59	0.75
insulin_Up	km	0.13	0.72	0.48
	rank	0.57	0.45	1.15
med_group_High	km	86.85	<0.005	66.21
	rank	128.33	<0.005	96.41
med_group_Moderate	km	5.66	0.02	5.84
	rank	17.48	<0.005	15.07
med_group_Very High	km	125.86	<0.005	94.62
	rank	80.55	<0.005	61.61

1. Variable 'med_group_Moderate' failed the non-proportional test: p-value is <5e-05.
- Advice: with so few unique values (only 2), you can include `strata=['med_group_Moderate', ...]` in the call in `.fit`. See documentation in link [E] below.
2. Variable 'med_group_High' failed the non-proportional test: p-value is <5e-05.
- Advice: with so few unique values (only 2), you can include `strata=['med_group_High', ...]` in the call in `.fit`. See documentation in link [E] below.
3. Variable 'med_group_Very High' failed the non-proportional test: p-value is <5e-05.
- Advice: with so few unique values (only 2), you can include `strata=['med_group_Very High', ...]` in the call in `.fit`. See documentation in link [E] below.

[A] https://lifelines.readthedocs.io/en/latest/jupyter_notebooks/Proportional%20hazard%20assumption.html
[B] https://lifelines.readthedocs.io/en/latest/jupyter_notebooks/Proportional%20hazard%20assumption.html#Bin-variable-and-stratify-on-it
[C] https://lifelines.readthedocs.io/en/latest/jupyter_notebooks/Proportional%20hazard%20assumption.html#Introduce-time-varying-covariates
[D] https://lifelines.readthedocs.io/en/latest/jupyter_notebooks/Proportional%20hazard%20assumption.html#Modify-the-functional-form
[E] https://lifelines.readthedocs.io/en/latest/jupyter_notebooks/Proportional%20hazard%20assumption.html#Stratification

Out [28]: []

Cox Proportional Hazards Model with Stratification

```
In [29]: # To group 'med_group' by num_medications
def group_meds(n):
    if n <= 5:
        return 'Low'
    elif n <= 10:
        return 'Moderate'
    elif n <= 15:
        return 'High'
    else:
        return 'Very High'
data['med_group'] = data['num_medications'].apply(group_meds)
```

```
In [30]: # To encode age and insulin (but NOT med_group)
data_model = pd.get_dummies(data[['age', 'insulin', 'duration', 'event']],
                             columns=['age', 'insulin'], drop_first=True)
```

```
In [31]: # To add unencoded med_group from original data
data_model['med_group'] = data['med_group']
```

```
In [32]: # To know the unique variables in med_group
print(data_model['med_group'].unique())
```

['Very High' 'Moderate' 'High' 'Low']

```
In [33]: # To fit stratified Cox model
cph = CoxPHFitter()
cph.fit(data_model, duration_col='duration', event_col='event', strata=['med_group'])
cph.print_summary()
```

model	lifelines.CoxPHFitter											
duration col	'duration'											
event col	'event'											
strata	med_group											
baseline estimation	breslow											
number of observations	54383											
number of events observed	6601											
partial log-likelihood	-59280.96											
time fit was run	2025-05-26 11:10:25 UTC											
	coef	exp(coef)	se(coef)	coef lower 95%	coef upper 95%	exp(coef) lower 95%	exp(coef) upper 95%	cmp to	z	p	-log2(p)	
age_[10-20)	0.84	2.31	0.60	-0.34	2.02	0.71	7.51	0.00	1.39	0.16	2.61	
age_[20-30)	1.79	6.02	0.58	0.65	2.94	1.92	18.86	0.00	3.08	<0.005	8.91	
age_[30-40)	1.49	4.46	0.58	0.35	2.63	1.43	13.94	0.00	2.57	0.01	6.62	
age_[40-50)	1.41	4.10	0.58	0.27	2.55	1.32	12.79	0.00	2.43	0.01	6.06	
age_[50-60)	1.27	3.56	0.58	0.13	2.41	1.14	11.08	0.00	2.19	0.03	5.12	
age_[60-70)	1.33	3.77	0.58	0.19	2.46	1.21	11.75	0.00	2.29	0.02	5.50	
age_[70-80)	1.30	3.67	0.58	0.16	2.44	1.18	11.42	0.00	2.24	0.02	5.32	
age_[80-90)	1.30	3.65	0.58	0.16	2.43	1.17	11.38	0.00	2.23	0.03	5.30	
age_[90-100)	1.29	3.64	0.58	0.15	2.44	1.16	11.43	0.00	2.21	0.03	5.22	

	coef	exp(coef)	se(coef)	coef lower 95%	coef upper 95%	exp(coef) lower 95%	exp(coef) upper 95%	cmp to	z	p	-log2(p)
insulin_Steady	-0.12	0.89	0.03	-0.17	-0.06	0.84	0.94	0.00	-3.86	<0.005	13.10
insulin_Up	-0.16	0.85	0.04	-0.23	-0.09	0.80	0.92	0.00	-4.43	<0.005	16.68

Concordance	0.53
Partial AIC	118583.91
log-likelihood ratio test	97.94 on 11 df
-log2(p) of ll-ratio test	50.96

```
In [34]: # To check PH assumptions
cph.check_assumptions(data_model, p_value_threshold=0.05)
```

Proportional hazard assumption looks okay.

Out [34]: []

```
In [35]: # To fit stratified Cox Model
cph.print_summary()
```

model	lifelines.CoxPHFitter										
duration col	'duration'										
event col	'event'										
strata	med_group										
baseline estimation	breslow										
number of observations	54383										
number of events observed	6601										
partial log-likelihood	-59280.96										
time fit was run	2025-05-26 11:10:25 UTC										
	coef	exp(coef)	se(coef)	coef lower 95%	coef upper 95%	exp(coef) lower 95%	exp(coef) upper 95%	cmp to	z	p	-log2(p)
age_[10-20]	0.84	2.31	0.60	-0.34	2.02	0.71	7.51	0.00	1.39	0.16	2.61
age_[20-30]	1.79	6.02	0.58	0.65	2.94	1.92	18.86	0.00	3.08	<0.005	8.91
age_[30-40]	1.49	4.46	0.58	0.35	2.63	1.43	13.94	0.00	2.57	0.01	6.62
age_[40-50]	1.41	4.10	0.58	0.27	2.55	1.32	12.79	0.00	2.43	0.01	6.06
age_[50-60]	1.27	3.56	0.58	0.13	2.41	1.14	11.08	0.00	2.19	0.03	5.12
age_[60-70]	1.33	3.77	0.58	0.19	2.46	1.21	11.75	0.00	2.29	0.02	5.50
age_[70-80]	1.30	3.67	0.58	0.16	2.44	1.18	11.42	0.00	2.24	0.02	5.32
age_[80-90]	1.30	3.65	0.58	0.16	2.43	1.17	11.38	0.00	2.23	0.03	5.30
age_[90-100]	1.29	3.64	0.58	0.15	2.44	1.16	11.43	0.00	2.21	0.03	5.22
insulin_Steady	-0.12	0.89	0.03	-0.17	-0.06	0.84	0.94	0.00	-3.86	<0.005	13.10
insulin_Up	-0.16	0.85	0.04	-0.23	-0.09	0.80	0.92	0.00	-4.43	<0.005	16.68

Concordance	0.53
Partial AIC	118583.91
log-likelihood ratio test	97.94 on 11 df
-log2(p) of ll-ratio test	50.96

```
In [ ]: 
```